Discovering Nulls and Outliers

DS-1004 Big Data

Advisor: Prof. Juliana Freire

MINIMIZERS

Diogo Mesquita Mihir Rana Kenil Tanna

https://github.com/ranamihir/big_data_project

INTRODUCTION

- Problem Statement:
  1. Null Value Detection
  2. Outlier Detection
     - Univariate outliers
     - Multivariate outliers
- Data Set Collection:
  - NYC Open Data
  - 50 data sets

PROBLEM FORMULATION

1. Data Cleaning
   - "$1.99" → 1.99, "1,000" → 1000, 10003 → "10003" (zip code)
2. Missing Value Treatment
   - "None", "N/A", " ", "-", "-999", "999", etc.
3. Outliers

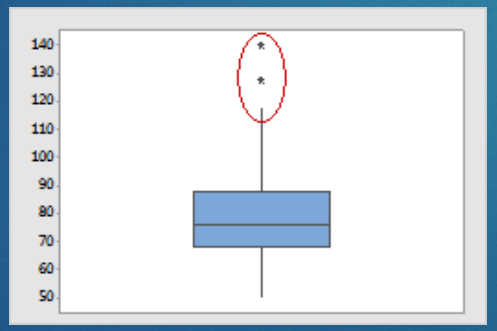

Source: link

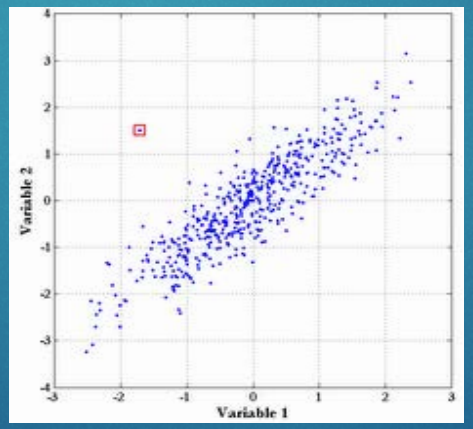

Source: link

METHODOLOGY

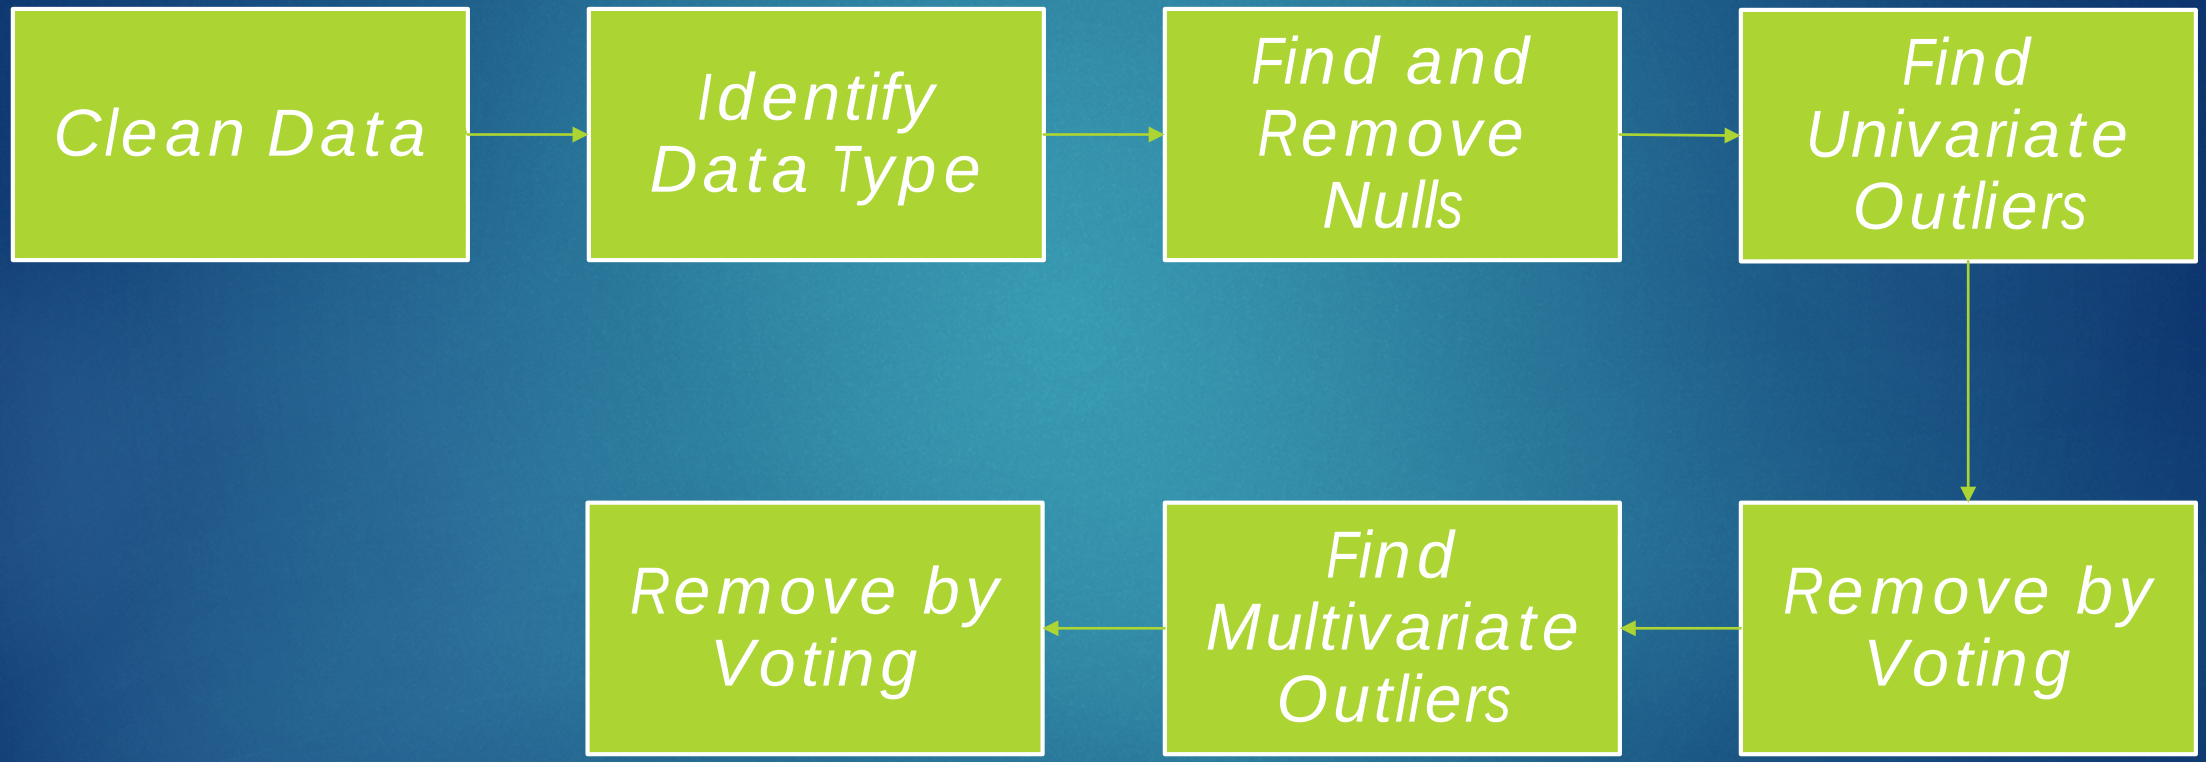

OUTLIER DETECTION

- *Nearest Neighbor based*
  - *DBSCAN*
- *Clustering based*
  - *k-Means*
- *Mixture of Parametric Distributions*
  - *Gaussian Mixture Models*

OUTLIER DETECTION

- Non-Parametric
  - Histogram / frequency – based
- Statistical Anomaly based
  - Box plot Rule
  - Gaussian model based (z-score)
  - Probabilistic models

KEY STRENGTHS

- End-to-end automated framework
- Box plot Rule at core
  - No input specific to particular data set / column required
- Robust
  - Multiple techniques optimizing different metrics
  - Voting / Intersection of multiple similar techniques
- Efficient
  - Remove univariate outliers before finding multivariate ones

RESULTS

- Data Cleaning
  - "$1.99“ → 1.99, "1,000“ → 1000, 10003 → “10003” (zipcode)
- Missing Value Treatment
- Outliers

REFERENCES

1. *Jason Brownlee. 2016. How To Handle Missing Values In Machine Learning Data With Weka. (Jun 2016). https://machinelearningmastery.com/ how-to-handle-missing-values-in-machine-learning-data-with-weka/*
2. *Varun Chandola, Arindam Banerjee, and Vipin Kumar. 2007. Anomaly Detection: A Survey. (2007).*
3. *Google. 2018. Locate Outliers, Google Cloud Dataprep Documentation, Google Cloud. (2018). https://cloud.google.com/dataprep/docs/html/Locate-Outliers_ 57344572*
4. *Ming Hua and Jian Pei. 2007. Cleaning Disguised Missing Data: A Heuristic Approach. In Proceedings of the 13th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '07). ACM, New York, NY, USA, 950–958. https://doi.org/10.1145/1281192.1281294*
5. *Ming Hua and Jian Pei. 2008. DiMaC: a disguised missing data cleaning tool. In KDD.*
6. *Sridhar Ramaswamy, Rajeev Rastogi, and Kyuseok Shim. 2000. Efficient Algorithms for Mining Outliers from Large Data Sets. SIGMOD Rec. 29, 2 (May 2000), 427–438. https://doi.org/10.1145/335191.335437*

REFERENCES

7. *Sheng Li, Ming Shao, and Yun Fu. 2015. Multi-view low-rank analysis for outlier detection. In Proceedings of the 2015 SIAM International Conference on Data Mining. SIAM, 748–756.*
8. *Sridhar Ramaswamy, Rajeev Rastogi, and Kyuseok Shim. 2000. Efficient Algorithms for Mining Outliers from Large Data Sets. In Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data (SIGMOD '00). ACM, New York, NY, USA, 427–438. https://doi.org/10.1145/342009.335437*
9. *Erich Schubert, Arthur Zimek, and Hans-Peter Kriegel. 2014. Local outlier detection reconsidered: a generalized view on locality with applications to spatial, video, and network outlier detection. Data Mining and Knowledge Discovery 28, 1 (01 Jan 2014), 190–237. https://doi.org/10.1007/s10618-012-0300-z*
10. *Arthur Zimek, Matthew Gaudet, Ricardo J.G.B. Campello, and Jörg Sander. 2013. Subsampling for Efficient and Effective Unsupervised Outlier Detection Ensembles. In Proceedings of the 19th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '13). ACM, New York, NY, USA, 428–436. https://doi.org/10.1145/2487575.2487676*

THANK YOU!

MINIMIZERS

Diogo Mesquita

Mihir Rana

Kenil Tanna